

A Physics 416 Fortran Tutorial

(based on the information at: <http://macams1.bo.infn.it/tutorial/>)

1. [What is Fortran?](#)
2. [Fortran basics](#)
3. [Variables, declarations, and types](#)
4. [Expressions and assignment](#)
5. [Logical expressions](#)
6. [The if statements](#)
7. [Loops](#)
8. [Arrays](#)
9. [Subprograms](#)
10. [Random numbers and Monte Carlo simulations](#)
11. [Simple input and output](#)
12. [Format statements](#)
13. [File I/O](#)
14. [Common blocks](#)
15. [data and block data](#)
16. [Debugging](#)
17. [Running Fortran on the Physics Department's VAX \(OHSTPY\) computer](#)
18. [A sample Fortran program for Lab 1](#)

1. What is Fortran?

Fortran is a general purpose programming language, mainly intended for mathematical computations in science applications (e.g. physics). Fortran is an acronym for FORMula TRANslation, and was originally capitalized as FORTRAN. However, following the current trend to only capitalize the first letter in acronyms, we will call it Fortran. Fortran was the first high-level programming language. The work on Fortran started in the 1950's at IBM and there have been many versions since. By convention, a Fortran version is denoted by the last two digits of the year the standard was proposed. Thus we have Fortran 66, Fortran 77 and Fortran 90 (95).

The most common Fortran version today is still Fortran 77, although Fortran 90 is growing in popularity. Fortran 95 is a revised version of Fortran 90 which is expected to be approved by ANSI soon (1996). There are also several versions of Fortran aimed at parallel computers. The most important one is High Performance Fortran (HPF), which is a de-facto standard.

Users should be aware that most Fortran 77 compilers allow a superset of Fortran 77, i.e. they allow non-standard extensions. In this tutorial we will emphasize standard ANSI Fortran 77.

Why learn Fortran?

Fortran is the dominant programming language used in scientific applications. It is therefore important for physics (or engineering) students to be able to read and modify Fortran code. From time to time, so-called experts predict that Fortran will rapidly fade in popularity and soon become extinct. This may actually happen as C (or C++) is rapidly growing in popularity. However, previous predictions of the downfall of Fortran have always been wrong. Fortran is the most enduring computer programming language in history. One of the main reasons Fortran has survived and will survive is *software inertia*. Once a company has spent many people-years and perhaps millions of dollars on a software product, it is unlikely to try to translate the software to a different language. Reliable software translation is a very difficult task and there's 40 years of Fortran code to replace!

Portability

A major advantage Fortran has is that it is standardized by ANSI (American National Standards Institute) and ISO (International Standards Organization). Consequently, if your program is written in ANSI Fortran 77 then it will run on any computer that has a Fortran 77 compiler. Thus, Fortran programs are portable across computer platforms

2. Fortran 77 Basics

A Fortran program is just a sequence of lines of text. The text has to follow a certain *syntax* to be a valid Fortran program. We start by looking at a simple example:

```
      program circle
      real r, area
c This program reads a real number r and prints
c the area of a circle with radius r.
      write (*,*) 'Give radius r:'
      read  (*,*) r
      area = 3.14159*r*r
      write (*,*) 'Area = ', area
      stop
      end
```

The lines that begin with a "c" are *comments* and have no purpose other than to make the program more readable for humans. Originally, all Fortran programs had to be written in all upper-case letters. Most people now write lower-case since this is more legible.

Program organization

A Fortran program generally consists of a main program (or driver) and possibly several subprograms (or procedures or subroutines). For now we will assume all the statements are in the main program; subprograms will be treated later. The structure of a main program is:

```
program name
  declarations
  statements
stop
end
```

In this tutorial, words that are in *italics* should not be taken as literal text, but rather as a generic description. The `stop` statement is optional and may seem superfluous since the program will stop when it reaches the end anyway but it is recommended to always terminate a program with the `stop` statement to emphasize that the execution flow stops there.

Column position rules

Fortran 77 is *not* a free-format language, but has a very strict set of rules for how the source code should be formatted. The most important rules are the column position rules:

```
Col. 1      : Blank, or a "c" or "*" for comments
Col. 2-5     : Statement label (optional)
Col. 6       : Continuation of previous line (optional)
Col. 7-72    : Statements
Col. 73-80   : Sequence number (optional, rarely used today)
```

Most lines in a Fortran 77 program starts with 6 blanks and ends before column 72, i.e. only the statement field is used. Note that Fortran 90 allows free format.

Comments

A line that begins with the letter "c" or an asterisk in the first column is a comment. Comments may appear anywhere in the program. Well-written comments are crucial to program readability. Commercial Fortran codes often contain about 50% comments. You may also encounter Fortran programs that use the exclamation mark (!) for comments. This is highly non-standard in Fortran 77, but is allowed in Fortran 90. The exclamation mark may appear anywhere on a line (except in positions 2-6).

Continuation

Occasionally, a statement does not fit into one single line. One can then break the statement into two or more lines, and use the continuation mark in position 6. Example:

```
c23456789 (This demonstrates column position!)

c The next statement goes over two physical lines
  area = 3.14159265358979
+      * r * r
```

Any character can be used instead of the plus sign as a continuation character. It is considered good programming style to use either the plus sign, an ampersand, or numbers (2 for the second line, 3 for the third, and so on).

Blank spaces

Blank spaces are ignored in Fortran 77. So if you remove all blanks in a Fortran 77 program, the program is still syntactically correct but almost unreadable for humans.

3. Variables, types, and declarations

Variable names

Variable names in Fortran consist of 1-6 characters chosen from the letters a-z and the digits 0-9. The first character must be a letter. (Note: Fortran 90 allows variable names of arbitrary length). Fortran 77 does not distinguish between upper and lower case, in fact, it assumes all input is upper case. However, nearly all Fortran 77 compilers will accept lower case. If you should ever encounter a Fortran 77 compiler that insists on upper case it is usually easy to convert the source code to all upper case.

Types and declarations

Every variable *should* be defined in a *declaration*. This establishes the *type* of the variable. The most common declarations are:

```
integer    list of variables
real       list of variables
double precision  list of variables
complex    list of variables
logical    list of variables
character  list of variables
```

The list of variables should consist of variable names separated by commas. Each variable should be declared exactly once. If a variable is undeclared, Fortran 77 uses a set of *implicit rules* to establish the type. This means all variables starting with the letters i-n are integers and all others are real. Many old Fortran 77 programs use these implicit rules, but *you should not!* The probability of errors in your program grows dramatically if you do not consistently declare your variables.

Integers and floating point variables

Fortran 77 has only one type for integer variables. Integers are usually stored as 32 bits (4 bytes) variables. Therefore, all integer variables should take on values in the range $[-m, m]$ where m is approximately 2×10^9 .

Fortran 77 has two different types for floating point variables, called `real` and `double precision`. While `real` is often adequate, some numerical calculations need very high precision and `double precision` should be used. Usually a `real` is a 4 byte variable and the `double precision` is 8 bytes, but this is machine dependent. Some non-standard Fortran versions use the syntax `real*8` to denote 8 byte floating point variables.

The parameter statement

Some constants appear many times in a program. It is then often desirable to define them only once, in the beginning of the program. This is what the `parameter` statement is for. It also makes programs more readable. For example, the circle area program should have been written like this:

```
program circle
  real r, area, pi
  parameter (pi = 3.14159)
c This program reads a real number r and prints
c the area of a circle with radius r.
  write (*,*) 'Give radius r:'
  read  (*,*) r
  area = pi*r*r
  write (*,*) 'Area = ', area
  stop
end
```

The syntax of the `parameter` statement is

parameter (*name = constant, ... , name = constant*)

The rules for the parameter statement are:

The "variable" defined in the parameter statement is not a variable but rather a constant whose value can never change

A "variable" can appear in at most one parameter statement

The parameter statement(s) must come before the first executable statement

Some good reasons to use the parameter statement are:

it helps reduce the number of typos

it is easy to change a constant that appears many times in a program

4. Expressions and assignment

Constants

The simplest form of an expression is a *constant*. There are 6 types of constants, corresponding to the 6 data types. Here are some integer constants:

1
0
-100
32767
+15

Then we have real constants:

1.0
-0.25
2.0E6
3.333E-1

The E-notation means that you should multiply the constant by 10 raised to the power following the "E". Hence, 2.0E6 is two million, while 3.333E-1 is approximately one third.

For constants that are larger than the largest real allowed, or that requires high precision, double precision should be used. The notation is the same as for real constants except the "E" is replaced by a "D".

Examples:

2.0D-1
1D99

Here 2.0D-1 is a double precision one-fifth, while 1D99 is a one followed by 99 zeros.

The next type is complex constants. This is designated by a pair of constants (integer or real), separated by a comma and enclosed in parentheses. Examples are:

(2, -3)
(1., 9.9E-1)

The first number denotes the real part and the second the imaginary part.

The fifth type is logical constants. These can only have one of two values:

.TRUE.
.FALSE.

Note that the dots enclosing the letters are required.

The last type is character constants. These are most often used as an *array* of characters, called a *string*. These consist of an arbitrary sequence of characters enclosed in apostrophes (single quotes):

```
'ABC'
'Anything goes!'
'It is a nice day'
```

Strings and character constants are case sensitive. A problem arises if you want to have an apostrophe in the string itself. In this case, you should double the apostrophe:

```
'It''s a nice day'
```

Expressions

The simplest expressions are of the form

operand operator operand

and an example is

```
x + y
```

The result of an expression is itself an operand, hence we can nest expressions together like

```
x + 2 * y
```

This raises the question of precedence: Does the last expression mean $x + (2*y)$ or $(x+2)*y$? The precedence of arithmetic operators in Fortran 77 are (from highest to lowest):

```
**      {exponentiation}
*,/     {multiplication, division}
+,-     {addition, subtraction}
```

All these operators are calculated left-to-right, except the exponentiation operator `**`, which has right-to-left precedence. If you want to change the default evaluation order, you can use parentheses.

The above operators are all binary operators. there is also the unary operator `-` for negation, which takes precedence over the others. Hence an expression like `-x+y` means what you would expect.

Extreme caution must be taken when using the division operator, which has a quite different meaning for integers and reals. If the operands are both integers, an integer division is performed, otherwise a real arithmetic division is performed. For example, `3/2` equals 1, while `3./2.` equals 1.5.

Assignment

The assignment has the form

```
variable_name = expression
```

The interpretation is as follows: Evaluate the right hand side and assign the resulting value to the variable on the left. The expression on the right may contain other variables, but these never change value! For example,

```
area = pi * r**2
```

does not change the value of `pi` or `r`, only `area`.

Type conversion

When different data types occur in the same expression, *type conversion* has to take place, either explicitly or implicitly. Fortran will do some type conversion implicitly. For example,

```
real x
x = x + 1
```

will convert the integer one to the real number one, and has the desired effect of incrementing `x` by one. However, in more complicated expressions, it is good programming practice to force the necessary type conversions explicitly. For numbers, the following functions are available:

```
int
real
```

```
db1e
ichar
char
```

The first three have the obvious meaning. `ichar` takes a character and converts it to an integer, while `char` does exactly the opposite.

Example: How to multiply two real variables `x` and `y` using double precision and store the result in the double precision variable `w`:

```
w = dble(x)*dble(y)
```

Note that this is different from

```
w = dble(x*y)
```

5. Logical expressions

Logical expressions can only have the value `.TRUE.` or `.FALSE.`. A logical expression can be formed by comparing arithmetic expressions using the following *relational operators*:

<code>.LT.</code>	means	less than (<code><</code>)
<code>.LE.</code>		less than or equal (<code><=</code>)
<code>.GT.</code>		greater than (<code>></code>)
<code>.GE.</code>		greater than or equal (<code>>=</code>)
<code>.EQ.</code>		equal (<code>=</code>)
<code>.NE.</code>		not equal (<code>/=</code>)

So you *cannot* use symbols like `<` or `=` for comparisons in Fortran 77.

For example: `(x.eq.y)` is valid while `(x=y)` is not valid in Fortran 77.

Logical expressions can be combined by the *logical operators* `.AND.` `.OR.` `.NOT.` which have the obvious meaning.

Logical variables and assignment

Truth values can be stored in *logical variables*. The assignment is analogous to the arithmetic assignment.

Example:

```
logical a, b
a = .TRUE.
b = a .AND. 3 .LT. 5/2
```

The order of precedence is important, as the last example shows. The rule is that arithmetic expressions are evaluated first, then relational operators, and finally logical operators. Hence `b` will be assigned `.FALSE.` in the example above.

Logical variables are seldom used in Fortran. But logical expressions are frequently used in conditional statements like the `if` statement.

6. The `if` statements

An important part of any programming language are the *conditional statements*. The most common such statement in Fortran is the `if` statement, which actually has several forms. The simplest one is the logical `if` statement:

```
if (logical expression) executable statement
```

This has to be written on one line. This example finds the absolute value of `x`:

```
if (x .LT. 0) x = -x
```

If more than one statement should be executed inside the `if`, then the following syntax should be used:

```
if (logical expression) then
  statements
endif
```

The most general form of the `if` statement has the following form:

```
if (logical expression) then
    statements
elseif (logical expression) then
    statements
:
:
else
    statements
endif
```

The execution flow is from top to bottom. The conditional expressions are evaluated in sequence until one is found to be true. Then the associated code is executed and the control jumps to the next statement after the `endif`.

Nested `if` statements

`if` statements can be nested in several levels. To ensure readability, it is important to use proper indentation. Here is an example:

```
if (x .GT. 0) then
    if (x .GE. y) then
        write(*,*) 'x is positive and x = y'
    else
        write(*,*) 'x is positive but x < y'
    endif
elseif (x .LT. 0) then
    write(*,*) 'x is negative'
else
    write(*,*) 'x is zero'
endif
```

You should avoid nesting many levels of `if` statements since things get hard to follow.

7. Loops

For repeated execution of similar things, *loops* are used. If you are familiar with other programming languages you have probably heard about *for*-loops, *while*-loops, and *until*-loops. Fortran 77 has only one loop construct, called the `do`-loop. The `do`-loop corresponds to what is known as a *for*-loop in other languages. Other loop constructs have to be simulated using the `if` and `goto` statements.

`do`-loops

The `do`-loop is used for simple counting. Here is a simple example that prints the cumulative sums of the integers from 1 through `n` (assume `n` has been assigned a value elsewhere):

```
integer i, n, sum
sum = 0
do 10 i = 1, n
    sum = sum + i
    write(*,*) 'i =', i
    write(*,*) 'sum =', sum
10 continue
```

The number 10 is a statement *label*. Typically, there will be many loops and other statements in a single program that require a statement label. The programmer is responsible for assigning a unique number to each label in each program (or subprogram). Recall that column positions 2-5 are reserved for statement labels. The numerical value of statement labels have no significance, so any integer numbers can be used. Typically, most programmers increment labels by 10 at a time.

The variable defined in the do-statement is incremented by 1 by default. However, you can define any other integer to be the *step*. This program segment prints the even numbers between 1 and 10 in decreasing order:

```
integer i

do 20 i = 10, 1, -2
  write(*,*) 'i =', i
20 continue
```

The general form of the do loop is as follows:

```
do label var = expr1, expr2, expr3
  statements
label continue
```

var is the loop variable (often called the *loop index*) which must be integer. *expr1* specifies the initial value of *var*, *expr2* is the terminating bound, and *expr3* is the increment (step).

Note: The do-loop variable must never be changed by other statements within the loop! This will cause great confusion.

Many Fortran 77 compilers allow do-loops to be closed by the enddo statement. The advantage of this is that the statement label can then be omitted since it is assumed that an enddo closes the nearest previous do statement. The enddo construct is widely used, but it is not a part of ANSI Fortran 77.

while-loops

The most intuitive way to write a while-loop is

```
while (logical expr) do
  statements
enddo
```

or alternatively,

```
do while (logical expr)
  statements
enddo
```

The statements in the body will be repeated as long as the condition in the while statement is true. Even though this syntax is accepted by many compilers, it is not ANSI Fortran 77. The correct way is to use if and goto:

```
label if (logical expr) then
  statements
  goto label
endif
```

Here is an example that calculates and prints all the powers of two that are less than or equal to 100:

```
integer n

n = 1
10 if (n .le. 100) then
  n = 2*n
  write (*,*) n
  goto 10
endif
```

until-loops

If the termination criterion is at the end instead of the beginning, it is often called an until-loop. The pseudocode looks like this:

```
do
    statements
until (logical expr)
```

Again, this should be implemented in Fortran 77 by using `if` and `goto`:

```
label continue
    statements
    if (logical expr) goto label
```

Note that the logical expression in the latter version should be the negation of the expression given in the pseudocode!

8. Arrays

Many scientific computations use vectors and matrices. The data type Fortran uses for representing such objects is the *array*. A one-dimensional array corresponds to a vector, while a two-dimensional array corresponds to a matrix. To fully understand how this works in Fortran 77, you will have to know not only the syntax for usage, but also how these objects are stored in memory in Fortran 77.

One-dimensional arrays

The simplest array is the one-dimensional array, which is just a linear sequence of elements stored consecutively in memory. For example, the declaration

```
real a(20)
```

declares `a` as a real array of length 20. That is, `a` consists of 20 real numbers stored contiguously in memory. By convention, Fortran arrays are indexed from 1 and up. Thus the first number in the array is denoted by `a(1)` and the last by `a(20)`. However, you may define an arbitrary index range for your arrays using the following syntax:

```
real b(0:19), weird(-162:237)
```

Here, `b` is exactly similar to `a` from the previous example, except the index runs from 0 through 19. `weird` is an array of length $237 - (-162) + 1 = 400$.

The type of an array element can be any of the basic data types. Examples:

```
integer i(10)
logical aa(0:1)
double precision x(100)
```

Each element of an array can be thought of as a separate variable. You reference the *i*'th element of array `a` by `a(i)`. Here is a code segment that stores the 10 first square numbers in the array `sq`:

```
integer i, sq(10)

do 100 i = 1, 10
    sq(i) = i**2
100 continue
```

A common bug in Fortran is that the program tries to access array elements that are out of bounds or undefined. This is the responsibility of the programmer, and the Fortran compiler will not detect any such bugs!

Two-dimensional arrays

Matrices are very important in linear algebra. Matrices are usually represented by two-dimensional arrays. For example, the declaration

```
real A(3,5)
```

defines a two-dimensional array of $3 \times 5 = 15$ real numbers. It is useful to think of the first index as the row index, and the second as the column index. Hence we get the graphical picture:

```
(1,1) (1,2) (1,3) (1,4) (1,5)
(2,1) (2,2) (2,3) (2,4) (2,5)
(3,1) (3,2) (3,3) (3,4) (3,5)
```

Two-dimensional arrays may also have indices in an arbitrary defined range. The general syntax for declarations is:

```
name (low_index1 : hi_index1, low_index2 : hi_index2)
```

The total size of the array is then

```
size = (hi_index1-low_index1+1)*(hi_index2-low_index2+1)
```

It is quite common in Fortran to declare arrays that are larger than the matrix we want to store. (This is because Fortran does not have dynamic storage allocation.) This is perfectly legal. Example:

```
real A(3,5)
integer i,j
c
c   We will only use the upper 3 by 3 part of this array.
c
do 20 j = 1, 3
  do 10 i = 1, 3
    a(i,j) = real(i)/real(j)
10  continue
20 continue
```

The elements in the submatrix $A(1:3,4:5)$ are undefined. Do not assume these elements are initialized to zero by the compiler (some compilers will do this, but not all).

Storage format for 2-dimensional arrays

Fortran stores higher dimensional arrays as a contiguous linear sequence of elements. It is important to know that 2-dimensional arrays are stored *by column*. So in the above example, array element (1,2) will follow element (3,1). Then follows the rest of the second column, thereafter the third column, and so on.

Consider again the example where we only use the upper 3 by 3 submatrix of the 3 by 5 array $A(3,5)$. The 9 interesting elements will then be stored in the first nine memory locations, while the last six are not used. This works out neatly because the *leading dimension* is the same for both the array and the matrix we store in the array. However, frequently the leading dimension of the array will be larger than the first dimension of the matrix. Then the matrix will *not* be stored contiguously in memory, even if the array is contiguous. For example, suppose the declaration was $A(5,3)$ instead. Then there would be two "unused" memory cells between the end of one column and the beginning of the next column (again we are assuming the matrix is 3 by 3).

This may seem complicated, but actually it is quite simple when you get used to it. If you are in doubt, it can be useful to look at how the *address* of an array element is computed. Each array will have some memory address assigned to the beginning of the array, that is element (1,1). The address of element (i,j) is then given by

$$\text{addr}[A(i,j)] = \text{addr}[A(1,1)] + (j-1)*\text{lda} + (i-1)$$

where *lda* is the leading (i.e. column) dimension of *A*. Note that *lda* is in general different from the actual matrix dimension. Many Fortran errors are caused by this, so it is very important you understand the distinction!

Multi-dimensional arrays

Fortran 77 allows arrays of up to seven dimensions. The syntax and storage format are analogous to the two-dimensional case, so we will not spend time on this.

The dimension statement

There is an alternate way to declare arrays in Fortran 77. The statements

```
real A, x
dimension x(50)
dimension A(10,20)
```

are equivalent to

```
real A(10,20), x(50)
```

This dimension statement is considered old-fashioned style today.

9. Subprograms

When a program is more than a few hundred lines long, it gets hard to follow. Fortran codes that solve real engineering problems often have tens of thousands of lines. The only way to handle such big codes, is to use a *modular* approach and split the program into many separate smaller units called *subprograms*.

A subprogram is a (small) piece of code that solves a well defined subproblem. In a large program, one often has to solve the same subproblems with many different data. Instead of replicating code, these tasks should be solved by subprograms. The same subprogram can be invoked many times with different input data.

Fortran has two different types of subprograms, called *functions* and *subroutines*.

Functions

Fortran functions are quite similar to mathematical functions: They both take a set of input arguments (parameters) and return a value of some type. In the preceding discussion we talked about *user defined* subprograms. Fortran 77 also has some *built-in* functions.

A simple example illustrates how to use a function:

```
x = cos(pi/3.0)
```

Here *cos* is the cosine function, so *x* will be assigned the value 0.5 (if *pi* has been correctly defined; Fortran 77 has no built-in constants). There are many built-in functions in Fortran 77. Some of the most common are:

<i>abs</i>	<i>absolute value</i>
<i>min</i>	<i>minimum value</i>
<i>max</i>	<i>maximum value</i>
<i>sqrt</i>	<i>square root</i>
<i>sin</i>	<i>sine</i>
<i>cos</i>	<i>cosine</i>
<i>tan</i>	<i>tangent</i>
<i>atan</i>	<i>arctangent</i>
<i>exp</i>	<i>exponential (natural)</i>
<i>log</i>	<i>logarithm (natural)</i>

In general, a function always has a *type*. Most of the built-in functions mentioned above, however, are *generic*. So in the example above, *pi* and *x* could be either of type *real* or *double precision*. The compiler would check the types and use the correct version of *cos* (*real* or *double precision*). Unfortunately, Fortran is not really a *polymorphic* language so in general you have to be careful to match the types of your variables and your functions!

Now we turn to the user-written functions. Consider the following problem: A meteorologist has studied the precipitation levels in the Bay Area and has come up with a model $r(m,t)$ where r is the amount of rain, m is the month, and t is a scalar parameter that depends on the location. Given the formula for r and the value of t , compute the annual rainfall.

The obvious way to solve the problem is to write a loop that runs over all the months and sums up the values of r . Since computing the value of r is an independent subproblem, it is convenient to implement it as a function. The following main program can be used:

```

      program rain
      real r, t, sum
      integer m

      read (*,*) t
      sum = 0.0
      do 10 m = 1, 12
         sum = sum + r(m, t)
10    continue
      write (*,*) 'Annual rainfall is ', sum, 'inches'

      stop
      end

```

In addition, the function r has to be defined as a Fortran function. The formula the meteorologist came up with was

$$r(m,t) = t/10 * (m^2 + 14m + 46) \text{ if this is positive}$$

$$r(m,t) = 0 \text{ otherwise}$$

The corresponding Fortran function is

```

      real function r(m,t)
      integer m
      real t
      r = 0.1*t * (m**2 + 14*m + 46)
      if (r .LT. 0) r = 0.0
      return
      end

```

We see that the structure of a function closely resembles that of the main program. The main differences are:

- Functions have a type. This type must also be declared in the calling program.
- The return value should be stored in a variable with the same name as the function.
- Functions are terminated by the *return* statement instead of *stop*.

To sum up, the general syntax of a Fortran 77 function is:

```

      type function name (list-of-variables)
      declarations
      statements
      return
      end

```

The function has to be declared with the correct type in the calling program unit. The function is then called by simply using the function name and listing the parameters in parenthesis.

Subroutines

A Fortran function can essentially only return one value. Often we want to return two or more values (or sometimes none!). For this purpose we use the subroutine construct. The syntax is as follows:

```

      subroutine name (list-of-arguments)
      declarations
      statements
      return
      end

```

Note that subroutines have no type and consequently should not (cannot) be declared in the calling program unit.

We give an example of a very simple subroutine. The purpose of the subroutine is to swap two integers.

```

      subroutine iswap (a, b)

```

```

        integer a, b
c Local variables
        integer tmp

        tmp = a
        a = b
        b = tmp

        return
    end

```

Note that there are two blocks of variable declarations here. First, we declare the input/output parameters, i.e. the variables that are common to both the caller and the callee. Afterwards, we declare the *local variables*, i.e. the variables that can only be used within this subprogram. We can use the same variable names in different subprograms and the compiler will know that they are different variables that just happen to have the same names.

Call-by-reference

Fortran 77 uses the so-called *call-by-reference* paradigm. This means that instead of just passing the values of the function/subroutine arguments (*call-by-value*), the memory address of the arguments (pointers) are passed instead. A small example should show the difference:

```

        program callex
        integer m, n
c
        m = 1
        n = 2
        call iswap(m, n)
        write(*,*) m, n
        stop
    end

```

The output from this program is "2 1", just as one would expect. However, if Fortran 77 had been using call-by-value then the output would have been "1 2", i.e. the variables m and n were unchanged! The reason for this is that only the values of m and n had been copied to the subroutine iswap, and even if a and b were swapped inside the subroutine the new values would not have been passed back to the main program.

In the above example, call-by-reference was exactly what we wanted. But you have to be careful about this when writing Fortran code, because it is easy to introduce undesired *side effects*. For example, sometimes it is tempting to use an input parameter in a subprogram as a local variable and change its value. You should *never* do this since the new value will then propagate back to the calling program with an unexpected value!

10. Random numbers and Monte Carlo simulations

All computer simulations (e.g. rolling the dice, tossing a coin) make use of a *function* that gives a random number uniformly between [0,1), i.e. includes 0, but not 1. In Fortran this function is `ran(seed)`, where seed is an integer variable used to generate ("seed") the sequence of random numbers. Below is a sample program that will generate 10 random numbers in the interval [0,1).

```

        program random
        integer seed, n
        seed=7654321
c seed should be set to a large odd integer according to the Fortran manual
        n = 10
        do n=1,10
            r=ran(seed)
            write(6,*) n, r
c could use a * instead of 6 in the write statement
        enddo
        stop
    end

```

The seed is used to generate the random numbers. If two programs (or the same program run twice) use the same seed, then they will get the same sequence of random numbers. Try running the above program twice! If you want a different sequence of random numbers every time you run your program then you must change your seed. One way to have your program change the seed for you is to use a function that gives a number that depends on the time of day. For example using the function `secnds(x)` gives the number of seconds (minus x) since midnight. So, as long as you don't re-run your program too quickly, or run it exactly the same time on two different days, this Fortran program will give a different set of random numbers every time it is run.

```

      program random
      integer seed,seed1, n
      real x
      seed1=7654321
      x=0.0
c seed should be set to a large odd integer according to the manual
c secnds(x) gives number of seconds-x elapsed since midnight
c the 2*int(secnds(x)) is always even (int=gives integer) so seed is always odd
      seed=seed1+2*int(secnds(x))
      n = 10
      do n=1,10
      r=ran(seed)
      write(6,*) n, r
c could use a * instead of 6 in the write statement
      enddo
      stop
      end

```

The random number generator function only gives numbers in the interval [0,1). Sometimes we want random numbers in a different interval, e.g. [-1,1). A simple transformation can be used to change intervals. For example, if we want a random number (x) in the interval [a,b) we can do so using:

$$x = (b-a) * \text{ran}(\text{seed}) - a$$

Thus for the interval [-1,1) we get: $x = 2 * \text{ran}(\text{seed}) - 1$.

In fact, we can take our set of numbers from the `ran(seed)` function which have a uniform probability distribution in the interval [0,1) and turn them into a set of numbers that look like they come from just about any probability distribution with any interval that one can imagine!

A few examples:

`dice=int(1+6*ran(seed))` This generates the roll of a 6 sided die.

`g=sqrt(-2*log(ran(seed)))*cos(2*pi*ran(seed))`

This generates a random number from a gaussian distribution with mean=0 and variance=1. We assume that pi is already initialized to 3.14159 in the program.

`t= -a*log(ran(seed))` This generates a random number from an exponential distribution with lifetime =a.

Being able to transform the random numbers obtained from `ran(seed)` into any probability distribution function we want is extremely useful and forms the basis of all computer simulations.

This type of simulation often goes by the name "Monte Carlo". Why Monte Carlo? In the pre-computer era a popular way to obtain a set of random numbers was to use a roulette wheel, just the type found in the Mediterranean city of Monte Carlo, famous for its gambling casino. This all happened in the late 1940's. If this technique had become popular in the late 1950's we'd probably be calling it Las Vegas!

11. Simple I/O

An important part of any computer program is to handle input and output. In our examples so far, we have already used the two most common Fortran constructs for this: `read` and `write`. Fortran I/O can be quite complicated, so we will only describe some simpler cases in this tutorial.

Read and write

Read is used for input, while write is used for output. A simple form is

```
read (unit no, format no) list-of-variables
write(unit no, format no) list-of-variables
```

The unit number can refer to either standard input, standard output, or a file. This will be described in later section. The format number refers to a label for a format statement, which will be described shortly.

It is possible to simplify these statements further by using asterisks (*) for some arguments, like we have done in all our examples so far. This is sometimes called *list directed* read/write.

```
read (*,*) list-of-variables
write(*,*) list-of-variables
```

The first statement will read values from the standard input and assign the values to the variables in the variable list, while the second one writes to the standard output.

Examples

Here is a code segment from a Fortran program:

```
integer m, n
real x, y
read(*,*) m, n
read(*,*) x, y
```

We give the input through standard input (possibly through a data file directed to standard input). A data file consists of *records* according to traditional Fortran terminology. In our example, each record contains a number (either integer or real). Records are separated by either blanks or commas. Hence a legal input to the program above would be:

```
-1 100
-1.0 1e+2
```

Or, we could add commas as separators:

```
-1, 100
-1.0, 1e+2
```

Note that Fortran 77 input is line sensitive, so it is important to have the right number of input elements (records) on each line. For example, if we gave the input all on one line as

```
-1, 100, -1.0, 1e+2
```

then m and n would be assigned the values -1 and 100 respectively, but the last two values would be discarded, leaving x and y undefined.

12. Format statements

So far we have only showed *free format* input/output. This uses a set of default rules for how to output values of different types (integers, reals, characters, etc.). Often the programmer wants to specify some particular input or output format, e.g. how many decimals in real numbers. For this purpose Fortran 77 has the *format* statement. The same format statements are used for both input and output.

Syntax

```
write(*, label) list-of-variables
label format format-code
```

A simple example demonstrates how this works. Say you have an integer variable you want to print in a field 4 characters wide and a real number you want to print in fixed point notation with 3 decimal places.

```
write(*, 900) i, x
900 format (I4,F8.3)
```


The format label 900 is chosen somewhat arbitrarily, but it is common practice to number format statements with higher numbers than the control flow labels. After the keyword `format` follows the format codes enclosed in parenthesis. The code I4 stands for an integer with width four, while F8.3 means that the number should be printed using fixed point notation with field width 8 and 3 decimal places.

The format statement may be located anywhere within the program unit. There are two programming styles: Either the format statement follows directly after the read/write statement, or all the format statements are grouped together at the end of the (sub-)program.

Common format codes

The most common format code letters are:

- A - text string
- D - double precision numbers, exponent notation
- E - real numbers, exponent notation
- F - real numbers, fixed point format
- I - integer
- X - horizontal skip (space)
- / - vertical skip (newline)

The format code F (and similarly D, E) has the general form `Fw.d` where *w* is an integer constant denoting the field width and *d* is an integer constant denoting the number of significant digits.

For integers only the field width is specified, so the syntax is `Iw`. Similarly, character strings can be specified as `Aw` but the field width is often dropped.

If a number or string does not fill up the entire field width, spaces will be added. Usually the text will be adjusted to the right, but the exact rules vary among the different format codes.

For horizontal spacing, the `nx` code is often used. This means *n* horizontal spaces. If *n* is omitted, *n=1* is assumed. For vertical spacing (newlines), use the code `/`. Each slash corresponds to one newline. Note that each read or write statement by default ends with a newline (here Fortran differs from C).

Some examples

This piece of Fortran code

```
x = 0.025
write(*,100) 'x=', x
100 format (A,F)
write(*,110) 'x=', x
110 format (A,F5.3)
write(*,120) 'x=', x
120 format (A,E)
write(*,130) 'x=', x
130 format (A,E8.1)
```

produces the following output when we run it:

```
x=      0.0250000
x=0.025
x=  0.2500000E-01
x=  0.3E-01
```

Note how blanks are automatically padded on the left and that the default field width for real numbers is usually 14. We see that Fortran 77 follows the rounding rule that digits 0-4 are rounded downwards while 5-9 is rounded upwards.

In this example each write statement used a different format statement. But it is perfectly fine to use the same format statement from many different write statements. In fact, this is one of the main advantages of using format statements. This feature is handy when you print tables for instance, and want each row to have the same format.

Format strings in read/write statements

Instead of specifying the format code in a separate format statement, one can give the format code in the read/write statement directly. For example, the statement

```
write (*,'(A, F8.3)') 'The answer is x = ', x
```

is equivalent to

```
write (*,990) 'The answer is x = ', x
990 format (A, F8.3)
```

Sometimes text strings are given in the format statements, e.g. the following version is also equivalent:

```
write (*,999) x
999 format ('The answer is x = ', F8.3)
```

Implicit loops and repeat counts

Now let us do a more complicated example. Say you have a two-dimensional array of integers and want to print the upper left 5 by 10 submatrix with 10 values each on 5 rows. Here is how:

```
do 10 i = 1, 5
  write(*,1000) (a(i,j), j=1,10)
10 continue
1000 format (I6)
```

We have an explicit do loop over the rows and an *implicit* loop over the column index j.

Often a format statement involves repetition, for example

```
950 format (2X, I3, 2X, I3, 2X, I3, 2X, I3)
```

There is a shorthand notation for this:

```
950 format (4(2X, I3))
```

It is also possible to allow repetition without explicitly stating how many times the format should be repeated. Suppose you have a vector where you want to print the first 50 elements, with ten elements on each line. Here is one way:

```
write(*,1010) (x(i), i=1,50)
1010 format (10I6)
```

The format statements says ten numbers should be printed. But in the write statement we try to print 50 numbers. So after the ten first numbers have been printed, the same format statement is automatically used for the next ten numbers and so on.

13. File I/O

So far we have assumed that the input/output has been to the standard input or the standard output. It is also possible to read or write from *files* which are stored on some external storage device, typically a disk (hard disk, floppy) or a tape. In Fortran each file is associated with a *unit number*, an integer between 1 and 99. Some unit numbers are reserved: 5 is standard input, 6 is standard output.

Opening and closing a file

Before you can use a file you have to *open* it. The command is

```
open (list-of-specifiers)
```

where the most common specifiers are:

```

[UNIT=]  u
IOSTAT=  ios
ERR=     err
FILE=    fname
STATUS=  sta
ACCESS=  acc
FORM=    frm
RECL=    rl

```

The unit number *u* is a number in the range 9-99 that denotes this file (the programmer may choose any number but he/she has to make sure it is unique).

ios is the I/O status identifier and should be an integer variable. Upon return, *ios* is zero if the statement was successful and returns a non-zero value otherwise.

err is a label which the program will jump to if there is an error.

fname is a character string denoting the file name.

sta is a character string that has to be either NEW, OLD or SCRATCH. It shows the prior status of the file. A scratch file is a file that is created and deleted when the file is closed (or the program ends).

acc must be either SEQUENTIAL or DIRECT. The default is SEQUENTIAL.

frm must be either FORMATTED or UNFORMATTED. The default is UNFORMATTED.

rl specifies the length of each record in a direct-access file.

For more details on these specifiers, see a good Fortran 77 book.

After a file has been opened, you can access it by read and write statements. When you are done with the file, it should be closed by the statement

```
close ([UNIT=]u[, IOSTAT=ios, ERR=err, STATUS=sta])
```

where, as usual, the parameters in brackets are optional.

Read and write revisited

The only necessary change from our previous simplified read/write statements, is that the unit number must be specified. But frequently one wants to add more specifiers. Here is how:

```
read ([UNIT=]u, [FMT=]fmt, IOSTAT=ios, ERR=err, END=s)
write([UNIT=]u, [FMT=]fmt, IOSTAT=ios, ERR=err, END=s)
```

where most of the specifiers have been described above. The END=s specifier defines which statement label the program jumps to if it reaches end-of-file.

Example

You are given a data file with xyz coordinates for a bunch of points. The number of points is given on the first line. The file name of the data file is *points.dat*. The format for each coordinate is known to be F10.4. Here is a short program that reads the data into 3 arrays x,y,z:

```

program inpdat
c
c  This program reads n points from a data file and stores them in
c  3 arrays x, y, z.
c
      integer nmax, u

```

```

        parameter (nmax=1000, u=20)
        real x(nmax), y(nmax), z(nmax)

c   Open the data file
        open (u, FILE='points.dat', STATUS='OLD')

c   Read the number of points
        read(u,*) n
        if (n.GT.nmax) then
            write(*,*) 'Error: n = ', n, 'is larger than nmax =', nmax
            goto 9999
        endif

c   Loop over the data points
        do 10 i= 1, n
            read(u,100) x(i), y(i), z(i)
        10 enddo
        100 format (3(F10.4))

c   Close the file
        close (u)

c   Now we should process the data somehow...
c   (missing part)

        9999 stop
        end

```

14. Common blocks

Fortran 77 has no *global* variables, i.e. variables that are shared among several program units (subroutines). The only way to pass information between subroutines we have seen so far, is to use the subroutine parameter list. Sometimes this is inconvenient, e.g. when many subroutines share a large set of parameters. In such cases one can use a *common block*. This is a way to specify that certain variables should be shared among certain subroutines. But in general, the use of common blocks should be minimized.

Example

Suppose you have two parameters alpha and beta that many of your subroutines need. The following example shows how it can be done using common blocks.

```

program main
    some declarations
    real alpha, beta
    common /coeff/ alpha, beta

    statements
    stop
    end

subroutine sub1 (some arguments)
    declarations of arguments
    real alpha, beta
    common /coeff/ alpha, beta

    statements
    return
    end

subroutine sub2 (some arguments)
    declarations of arguments

```

```

real alpha, beta
common /coeff/ alpha, beta

statements
return
end

```

Here we define a common block with the name `coeff`. The content of the common block is the two variables `alpha` and `beta`. A common block can contain as many variables as you like. They do not need to all have the same type. Every subroutine that wants to use any of the variables in the common block has to declare the whole block.

Note that in this example we could easily have avoided common blocks by passing `alpha` and `beta` as parameters (arguments). A good rule is to try to avoid common blocks if possible. However, there are a few rare cases where there is no other solution.

Syntax

```

common / name / list-of-variables

```

You should know that

The `common` statement should appear together with the variable declarations, before the executable statements.

Different common blocks must have different names (just like variables). A variable can belong to more than one common block.

The variables in a common block do not need to have the same names each place they occur (although it is a good idea to do so), but they must be listed in the same order and have the same type.

To illustrate this, look at the following continuation of our example:

```

subroutine sub3 (some arguments)
declarations of arguments
real a, b
common /coeff/ a, b

statements
return
end

```

This declaration is equivalent to the previous version that used `alpha` and `beta`. It is recommended that you always use the same variable names for the same common block to avoid confusion. Here is a dreadful example:

```

subroutine sub4 (some arguments)
declarations of arguments
real alpha, beta
common /coeff/ beta, alpha

statements
return
end

```

Now `alpha` is the `beta` from the main program and vice versa. If you see something like this, it is probably a mistake. Such bugs are very hard to find.

Arrays in common blocks

Common blocks can include arrays, too. But again, this is not recommended. The major reason is flexibility. An example shows why this is such a bad idea. Suppose we have the following declarations in the main program:

```

program main
integer nmax
parameter (nmax=20)
integer n
real A(nmax, nmax)
common /matrix/ A, n, nmax

```

This common block contains first all the elements of *A*, then the integers *n* and *nmax*. Now assume you want to use the matrix *A* in some subroutines. Then you have to include the same declarations in all these subroutines, e.g.

```

subroutine sub1 (...)
integer nmax
parameter (nmax=20)
integer n
real A(nmax, nmax)
common /matrix/ A, n, nmax

```

Arrays with variable dimensions cannot appear in common blocks, thus the value of *nmax* has to be exactly the same as in the main program. Recall that the size of a matrix has to be known at compile time, hence *nmax* has to be defined in a parameter statement. It would be tempting to delete the parameter statement in the subroutine since *nmax* belongs to the common block, but this would be illegal.

This example shows there is usually nothing to gain by putting arrays in common blocks. Hence the preferred method in Fortran 77 is to pass arrays as arguments to subroutines (along with the leading dimensions).

15. data and block data

The data statement

The data statement is another way to input data that are known at the time when the program is written. It is similar to the assignment statement. The syntax is:

```
data list-of-variables/ list-of-values/, ...
```

where the three dots means that this pattern can be repeated. Here is an example:

```
data m/10/, n/20/, x/2.5/, y/2.5/
```

We could also have written this

```
data m,n/10,20/, x,y/2*2.5/
```

We could have accomplished the same thing by the assignments

```

m = 10
n = 20
x = 2.5
y = 2.5

```

The data statement is more compact and therefore often more convenient. Notice especially the shorthand notation for assigning identical values repeatedly.

The data statement is performed only once, right before the execution of the program starts. For this reason, the data statement is mainly used in the main program and not in subroutines.

The data statement can also be used to initialize arrays (vectors, matrices). This example shows how to make sure a matrix is all zeros when the program starts:

```
real A(10,20)
```

```
data A/ 200 * 0.0/
```

Some compilers will automatically initialize arrays like this but not all, so if you rely on array elements to be zero it is a good idea to follow this example. Of course you can initialize arrays to other values than zero. You may even initialize individual elements:

```
data A(1,1)/ 12.5/, A(2,1)/ -33.3/, A(2,2)/ 1.0/
```

Or you can list all the elements for small arrays like this:

```
integer v(5)
real B(2,2)
data v/10,20,30,40,50/, B/1.0,-3.7,4.3,0.0/
```

The values for two-dimensional arrays will be assigned in column-first order as usual.

The block data statement

The data statement cannot be used for variables contained in a common block. There is a special "subroutine" for this purpose, called `block data`. It is not really a subroutine, but it looks a bit similar because it is given as a separate program unit. Here is an example:

```
block data
integer nmax
parameter (nmax=20)
real v(nmax), alpha, beta
common /vector/v,alpha,beta
data v/20*100.0/, alpha/3.14/, beta/2.71/
end
```

Just as the data statement, block data is executed once before the execution of the main program starts. The position of the block data "subroutine" in the source code is irrelevant (as long as it is not nested inside the main program or a subprogram).

16. Debugging hints

It has been estimated that about 90% of the time it takes to develop commercial software is spent debugging and testing. This shows how important it is to write good code in the first place.

Still, we all discover bugs from time to time. Here are some hints for how to track them down.

Useful compiler options

Most Fortran compilers will have a set of options you can turn on if you like. The following compiler options are specific to the Sun Fortran 77 compiler, but most compilers will have similar options (although the letters may be different).

-ansi This will warn you about all non-standard Fortran 77 extensions in your program.

-u This will override the implicit type rules and make all variables undeclared initially. If your compiler does not have such an option, you can achieve almost the same thing by adding the declaration `implicit none (a-z)` in the beginning of each (sub-)program.

-C Check for array bounds. The compiler will try to detect if you access array elements that are out of bounds. It cannot catch all such errors, however.

Some common errors

Here are some common errors to watch out for:

Make sure your lines end at column 72. The rest will be ignored!

Have you done integer division when you wanted real division?

Do your parameter lists in the calling and the called program match?

Do your common blocks match?

Debugging tools

If you have a bug, you have to try to locate it. Syntax errors are easy to find. The problem is when you have run-time errors. The old-fashioned way to find errors is to add *write* statements in your code and try to track the values of your variables. This is a bit tedious since you have to recompile your source code every time you change the program slightly. Today one can use special *debuggers* which is a convenient tool. You can step through a program line by line or define your own break points, you can display the values of the variables you want to monitor, and much more.

17. Running Fortran on the Physics Department's VAX (OHSTPY) Computer

=> *do not type the “ ”s, just words inside the “ ”s!*

- 1) Log into OHSTPY.
- 2) Create your program by typing "edit name.for" and type in the program. Here, *name* is whatever you wish to call your program. For example, test.for or lab1.for. The extension *.for* tells the computer that this file contains Fortran code. The *edit* command runs a program that allows you to create and edit a file (*name.for*) on the computer's disk. When done entering your Fortran commands, hit <CTRL>z and type "exit" to save the recent edits and quit or "quit" to quit without saving the edits.
- 3) Compile the program by typing "fortran name", where name could be *test* if your file was named *test.for*. The compiler is a computer program that reads your Fortran code and translates it into code that the computer actually runs. The compiler creates a new file called name.obj (e.g. *test.obj*). If there are any syntax errors in your Fortran code the compiler will flag them and print out an error message(s) to the computer screen. You must fix all of the errors in the program before you can run the program.
- 4) Link your Fortran code (e.g. *test.for*) with other programs necessary to run your program by typing "link test". This creates a new file (an "executable") called name.exe (e.g. *test.exe*). Once the executable file is created, execute (run) the program by typing "run name" (e.g. "run test").
- 5) Check your output carefully! Just because the program runs does not mean it actually works!

18. A sample Fortran program for Lab 1

```
C      This program simulates rolling a six sided dice, each side has 1 to 6 dots.
C      The number of times the die is tossed is given by the integer variable: roll.
C
C      The result of an individual toss (a 1 or 2,..or 6 dots)
C      is given by the integer variable: dice.
C
C      The real variable count is an ARRAY that stores six numbers.
C      Count keeps track of the number of times a 1 (2,etc) was rolled.
C      For example if array(1) = 5 and array(2) = 7 then a one was
C      rolled 5 times and a two was rolled 7 times.
C
C      FORTRAN is based on lines and columns, usually one instruction per line.
C234567
C      In FORTRAN computer instructions start in column 7.
C      In FORTRAN a C in column 1 flags this line as a comment.
C      In FORTRAN a character in column 6 flags this as a continuation line.
C      Next lines starts the actual FORTRAN code:
      implicit none
      character key
      real count(6), RAN
      integer i, dice, roll, seed
C
C      Special precaution in integer division: dividing one integer
C      by another integer gives another integer, e.g. i = 3, j = 2,
C      then c = i/j = 1.0. To get the correct answer, define i and
C      j as real numbers.
C
C      RAN is a computer function that gives us a random number in
C      the range [0,1]. Initialize the random number generator RAN
C      with a large odd integer seed.
C
      seed = 432211111
10    continue
C      "10" is the statement number of the continue command.
C      Statement numbers should start at column 2 or later and end at column 5.
C      All statement numbers must be a unique integer.
C
C      Initialization of variables
C      Here we have a do loop, which repeats itself 6 times
      do dice = 1, 6
         count(dice) = 0
      end do
C      Write to unit 6 (screen) and read from unit 5 (keyboard).
      write(6,*)'How many rolls of dice :'
      read(5,*)roll
C      To understand how you generate a value for the dice, let's pick a
C      random number, say 0.6:
      dice = INT(1 + 6*RAN(seed)) = INT(1 + 6*0.6) = INT(4.6) = 4
C      i.e. INT truncates off the decimal numbers to give an integer.
      do i = 1, roll
         dice = INT(1 + 6*RAN(seed))
         count(dice) = count(dice) + 1
      end do
      do i = 1, 6
         write(6,*)i, count(i)
      end do
C      The ">" character at 6th character location indicates
C      continuation of a program line.
C
```

```
    write(6,*) 'Hit <return> to continue or type q or Q',  
>      ' and <return> to quit the program.'  
    read(5,20)key  
C      Define key as a one character variable ("a1").  If you use  
C      * instead of 20, you do not explicitly specify the format.  
20     format(a1)  
       if (key.ne.'q'.and.key.ne.'Q') go to 10  
       end
```